Rishabh Maheshway

rf.rishabh@gmail.com | +91 − 8427119320 | Google Scholar | Website | GitHub | LinkedIn

Summary

I am Rishabh, and I work at the intersection of AI research and engineering, with experience across text, multimodal and multilingual domains. My work covers the post-training stack, from synthetic data and fine-tuning to reward modeling, agentic reinforcement learning, evaluations, and adversarial robustness.

Work Experience

- **ServiceNow** - *Senior Research Scientist* — **Hyderabad, India** — *Aug '23 – Present*
  I led and contributed to post-training efforts for large language models, spanning data curation, fine-tuning, reinforcement learning based alignment, and benchmarking. I also worked on implementing long horizon agentic training and evaluation setups in interactive environments. Selected works include: Apriel-15B Series | Datasets, Alignment and Benchmarking | Agentic RL.

- **Facebook AI Research** - *AI Resident* — **California, U.S.** — *Nov '21 – Dec '22*
  I conducted research to improve the robustness and text understanding capabilities in multimodal models, contributing to Selective VQA and end-to-end text detection (Advisor: Marcus Rohrbach).

- **Verisk AI** – *Research Intern* — **Hyderabad, India** — *May – Oct 2021*
  Worked on joint language and vision understanding of multimodal content and semantic understanding of natural language documents.

- **Google Summer of Code** – *Software Developer Intern* — **Remote, India** — *Apr – Sept 2018*
  Developed an application enabling users to report incidents, facilitating nearby assistance and communication.

Publications

1. **Learning What to Remember: Agentic Note Writing for Long-Horizon QA**
   First Author *Accepted at ServiceNow internal Now-AI Conference | Under Submission in ACL 2026.*

2. **INCLUDE: Evaluating Multilingual Language Understanding with Regional Knowledge.**
   Core Contributor. *International Conference on Learning Representations* (*ICLR 2025*).

3. **M-RewardBench: Evaluating Reward Models in Multilingual Settings.** Srishti Gureja*, Lester James V. Miranda*, Shayekh Bin Islam*, **Rishabh Maheshwary***, Drishti Sharma, Gusti Winata, Nathan Lambert, Sebastian Ruder, Sara Hooker, Marzieh Fadaee.
   Association for Computational Linguistics (*ACL 2025*).

4. **Variable Layerwise Quantization: A Simple and Effective Approach to Quantize LLMs.** Razvan-Gabriel Dumitru, Vikas Yadav, **Rishabh Maheshwary**, Paul Ioan Clotan, Sathwik Tejaswi Madhusudhan, Mihai Surdeanu.
   Association for Computational Linguistics (*ACL 2025*).

5. **M2Lingual: Enhancing Multilingual, Multi-Turn Instruction Alignment in LLMs.** **Rishabh Maheshwary**, Vikas Yadav, Hoang Nguyen, Khyati Mahajan, Sathwik Tejaswi Madhusudhan.
   North American Chapter of the Association for Computational Linguistics (*NAACL 2025*).

6. **Enhancing Alignment using Curriculum Learning & Ranked Preferences.** Pulkit Pattnaik*, **Rishabh Maheshwary***, Kelechi Ogueji, Vikas Yadav, Sathwik Tejaswi Madhusudhan.
Findings of Empirical Methods in Natural Language Processing (*EMNLP 2024*).

7. **Improving Selective Visual Question Answering by Learning from Your Peers.** Corentin Dancette, Spencer Whitehead, **Rishabh Maheshwary**, Ramakrishna Vedantam, Stefan Scherer, Xinlei Chen, Matthieu Cord, Marcus Rohrbach.
IEEE/CVF Conference on Computer Vision and Pattern Recognition (*CVPR 2023*).

8. **Practice Makes a Solver Perfect: Data Augmentation Methods for Math Word Problem Solvers.** Vivek Kumar, **Rishabh Maheshwary**, Vikram Pudi.
North American Chapter of the Association for Computational Linguistics (*NAACL*) *2022.*

9. **A Strong Baseline for Query Efficient Attacks in a Black Box Setting. Rishabh Maheshwary***, Saket Maheshwary*, Vikram Pudi.
Empirical Methods in Natural Language Processing (*EMNLP 2021*).

10. **Adversarial Examples for Evaluating Math Word Problem Solvers. Rishabh Maheshwary***, Vivek Kumar*, Vikram Pudi.
Findings of Empirical Methods in Natural Language Processing EMNLP 2021.

11. **Generating Natural Language Attacks in a Hard Label Black Box Setting. Rishabh Maheshwary**, Saket Maheshwary, Vikram Pudi.
Association for the Advancement of Artificial Intelligence (*AAAI 2021*).

12. **A Context Aware Approach for Generating Natural Language Attacks. Rishabh Maheshwary**, Saket Maheshwary, Vikram Pudi.
Association for the Advancement of Artificial Intelligence (*AAAI 2021*).

---

* Equal Contribution

Education

- **International Institute of Information Technology, Hyderabad** — **Hyderabad, India**
  MS by Research in Computer Science and Engineering | CGPA: 8.7/10 — *2019 – 2021*
- **University Institute of Engineering and Technology, Panjab University** — **Chandigarh, India**
  BTech in Computer Science and Engineering | CGPA: 8.4/10 — *2015 – 2019*

Miscellaneous

- Actively reviewing for ACL, TMLR, ICLR, ECCV, NeurIPS and CoNLL. — *2021 – Present*
- **Google Summer of Code** and **Google CodeIn** mentor. — *2018*
- Ranked 1^{st} out of 100+ teams in **CODETRIX** (National level coding contest). — *2017*
- Ranked 3^{rd} out of 100+ teams in **CODE-IT** (National level coding contest). — *2016*

Programming Languages and Technologies

- Python, PyTorch, C++, C, Shell, Git
- NLP, Multimodal vision & language, Reinforcement learning