

# Rishabh Maheshwary

rf.rishabh@gmail.com | [Google Scholar](#) | [Website](#) | [GitHub](#) | [LinkedIn](#)

---

## Summary

I am Rishabh, and I work at the intersection of AI research and engineering, with experience across text, multimodal and multilingual domains. My work covers the post-training stack, from synthetic data and fine-tuning to reward modeling, agentic reinforcement learning, evaluations, and adversarial robustness.

---

## Work Experience

- **Cartesia AI** - *Member of Technical Staff* **Bangalore, India**  
Oct '26 – Present  
I am working on post-training & evals for LLM & text-to-speech models.
  - **ServiceNow** - *Senior Research Scientist* **Hyderabad, India**  
Aug '23 – Sept '26  
I led & contributed to LLM post-training across data synthesis, fine-tuning, long-horizon agentic RL and evaluations. Some of the selected works are — [Datasets, Alignment and Benchmarking](#) | [Apriel-15B Series](#) | [Agentic RL](#)
  - **Facebook AI Research** - *AI Resident* **California, U.S.**  
Nov '21 – Dec '22  
I worked on improving the robustness and text understanding in multimodal LLMs. [Abstain VQA](#) & [end to end text detection](#) (Advisor: [Marcus Rohrbach](#)).
  - **Verisk AI** – *Research Intern* **Hyderabad, India**  
May – Oct 2021  
Worked on joint language and vision understanding of multimodal content and semantic understanding of natural language documents.
- 

## Publications

1. **R2VAgent — Teaching SLMs When to Ask for Help** Raghu Vamshi Hemadri\*, Humaira Firdowse Mohammed\*, **Rishabh Maheshwary**, Srivatsava Daruru, Sagar Davasam, Vikas Yadav, Srinivas Sunkara, Sai Rajeswar | *Under submission at ICLR 2027*.
2. **Do Enterprise Systems Need World Models? The Importance of Context to Infer Dynamics** Jishnu Sethumadhavan Nair\*, Patrice Bechard\*, **Rishabh Maheshwary\***, Surajit Dasgupta, Sravan Ramachandran, Aakash Bhagat, Shruthan Radhakrishna, Pulkit Pattnaik, Johan Obando-Ceron, Shiva Krishna Reddy Malay, Sagar Davasam, Seganrasan Subramanian, Vipul Mittal, Sridhar Krishna Nemala, Christopher Pal, Srinivas Sunkara, Sai Rajeswar | *NeurIPS 2026 (Main Conference)*.
3. **Augmenting LLM Reasoning with Dynamic Notes Writing for Complex QA** **Rishabh Maheshwary**, Masoud Hashemi, Khyati Mahajan, Shiva Krishna Reddy Malay, Sai Rajeswar, Sathwik Tejaswi Madhusudhan, Spandana Gella, Vikas Yadav | *LREC 2026 | ServiceNow internal Now-AI Conference*.
4. **INCLUDE: Evaluating Multilingual Language Understanding with Regional Knowledge.** **Core Contributor** | *ICLR 2025*
5. **M-RewardBench: Evaluating Reward Models in Multilingual Settings.** Srishti Gureja\*, Lester James V. Miranda\*, Shayekh Bin Islam\*, **Rishabh Maheshwary\***, Drishti Sharma, Gusti Winata, Nathan Lambert, Sebastian Ruder, Sara Hooker, Marzieh Fadaee. | *ACL 2025*.

6. **Variable Layerwise Quantization: A Simple and Effective Approach to Quantize LLMs.** Razvan-Gabriel Dumitru, Vikas Yadav, **Rishabh Maheshwary**, Paul Ioan Clotan, Sathwik Tejaswi Madhusudhan, Mihai Surdeanu | *ACL* 2025.
7. **M2Lingual: Enhancing Multilingual, Multi-Turn Instruction Alignment in LLMs.** **Rishabh Maheshwary**, Vikas Yadav, Hoang Nguyen, Khyati Mahajan, Sathwik Tejaswi Madhusudhan | *NAACL* 2025.
8. **Enhancing Alignment using Curriculum Learning & Ranked Preferences.** Pulkit Pattnaik\*, **Rishabh Maheshwary\***, Kelechi Ogueji, Vikas Yadav, Sathwik Tejaswi Madhusudhan. | *EMNLP* 2024.
9. **Improving Selective Visual Question Answering by Learning from Your Peers.** Corentin Dancette, Spencer Whitehead, **Rishabh Maheshwary**, Ramakrishna Vedantam, Stefan Scherer, Xinlei Chen, Matthieu Cord, Marcus Rohrbach. | *CVPR* 2023.
10. **Practice Makes a Solver Perfect: Data Augmentation Methods for Math Word Problem Solvers.** Vivek Kumar, **Rishabh Maheshwary**, Vikram Pudi. | *NAACL* 2022.
11. **A Strong Baseline for Query Efficient Attacks in a Black Box Setting.** **Rishabh Maheshwary\***, Saket Maheshwary\*, Vikram Pudi. | *EMNLP* 2021.
12. **Adversarial Examples for Evaluating Math Word Problem Solvers.** **Rishabh Maheshwary\***, Vivek Kumar\*, Vikram Pudi. | *Findings of Empirical Methods in Natural Language Processing EMNLP* 2021.
13. **Generating Natural Language Attacks in a Hard Label Black Box Setting.** **Rishabh Maheshwary**, Saket Maheshwary, Vikram Pudi. | *AAAI* 2021.
14. **A Context Aware Approach for Generating Natural Language Attacks.** **Rishabh Maheshwary**, Saket Maheshwary, Vikram Pudi. | *AAAI* 2021.

---

\* Equal Contribution

## Education

- |                                                                                                                                                   |                                         |
|---------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------|
| • <b>International Institute of Information Technology, Hyderabad</b><br><i>MS by Research in Computer Science and Engineering</i>   CGPA: 8.7/10 | <b>Hyderabad, India</b><br>2019 – 2021  |
| • <b>University Institute of Engineering and Technology, Panjab University</b><br><i>BTech in Computer Science and Engineering</i>   CGPA: 8.4/10 | <b>Chandigarh, India</b><br>2015 – 2019 |
- 

## Miscellaneous

- |                                                                                                |                |
|------------------------------------------------------------------------------------------------|----------------|
| • Actively reviewing for ACL, TMLR, ICLR, ECCV, NeurIPS and CoNLL.                             | 2021 – Present |
| • <b>Google Summer of Code</b> and <b>Google CodeIn</b> mentor.                                | 2018           |
| • Ranked 1 <sup>st</sup> out of 100+ teams in <b>CODETRIX</b> (National level coding contest). | 2017           |
| • Ranked 3 <sup>rd</sup> out of 100+ teams in <b>CODE-IT</b> (National level coding contest).  | 2016           |
- 

## Programming Languages and Technologies

- Python, PyTorch, C++, C, Shell, Git
- NLP, Multimodal vision & language, Reinforcement learning